

---

---

# The Ins & Outs of Model Inheritance

---

---

— Blythe J Dunham —

blythe@snowgiraffe.com

---

---

# Today's Adventure Includes

- Composition vs Inheritance
- Three Types of Model Inheritance
  - Abstract Models ([The Hum of the Giraffe](#))
  - Multi Table Inheritance ([Don't Get Eaten by the Lion](#))
  - Proxy Models and Single Table Inheritance  
([Adventures in Middle Earth](#))
- Two Alternatives to Model Inheritance
- Avoiding Model Inheritance



# Today's Adventure Includes

- Composition vs Inheritance
- Three Types of Model Inheritance
  - Abstract Models ([The Hum of the Giraffe](#))
  - Multi Table Inheritance ([Don't Get Eaten by the Lion](#))
  - Proxy Models and Single Table Inheritance  
([Adventures in Middle Earth](#))
- Two Alternatives to Model Inheritance
- Avoiding Model Inheritance



# “Favor Composition Over Inheritance”

Design Patterns: Elements of Reusable Object-Oriented Software, 1994

# Composition: “has-a”

- A giraffe has a (blue) tongue
- A hobbit has hairy feet
- A lion has big mouth full of teeth



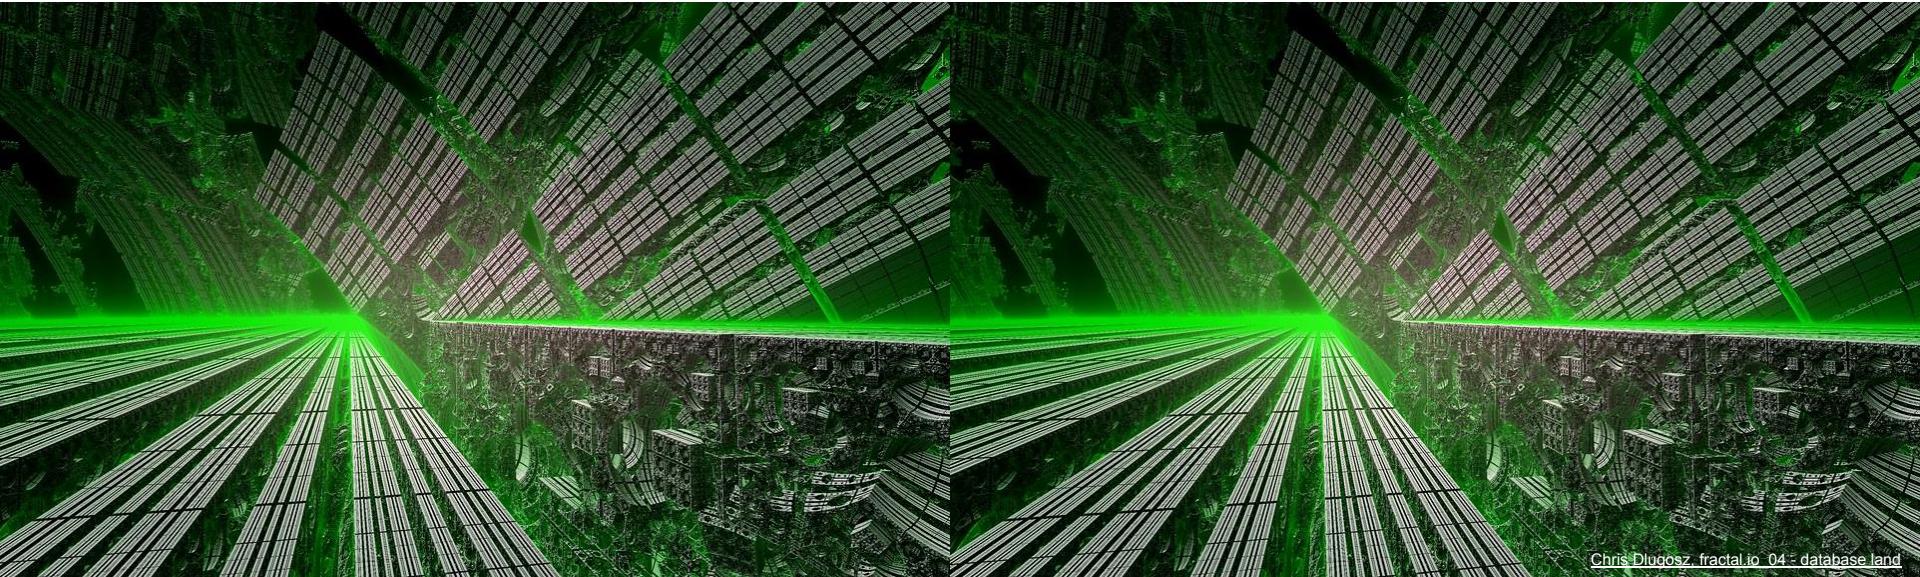
# Inheritance: “is-a”



- A giraffe is a (glorious) animal
- A lion is a big cat
- An elf is an immortal being

# Relational Databases Don't Implement Inheritance

Because most RDBS aren't object-oriented they do not support inheritance.  
While composition has a natural mapping, inheritance is not intuitive.

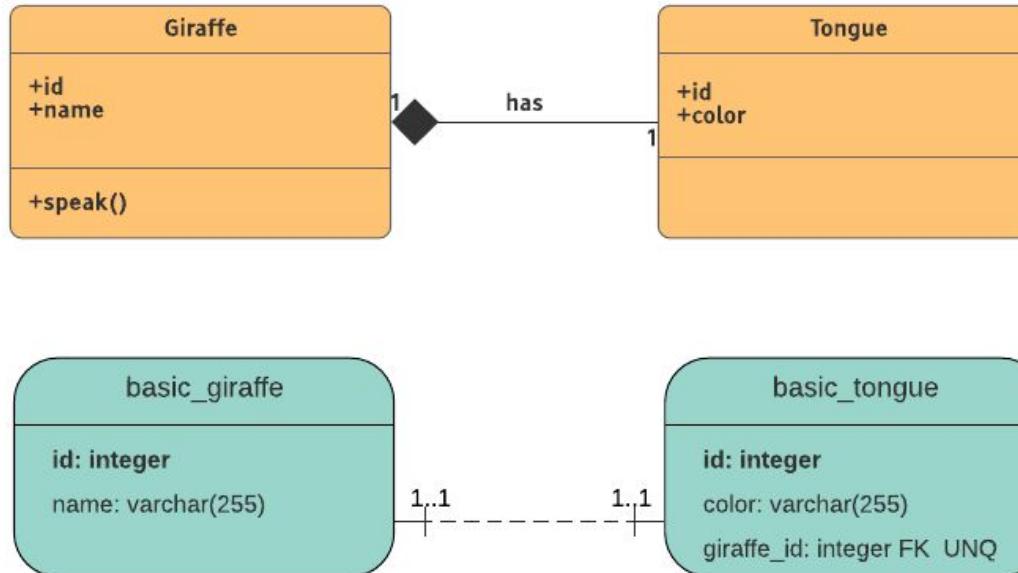


# Django 101: Modeling Composition

```
class Giraffe(models.Model):
    name = models.CharField(max_length=255)

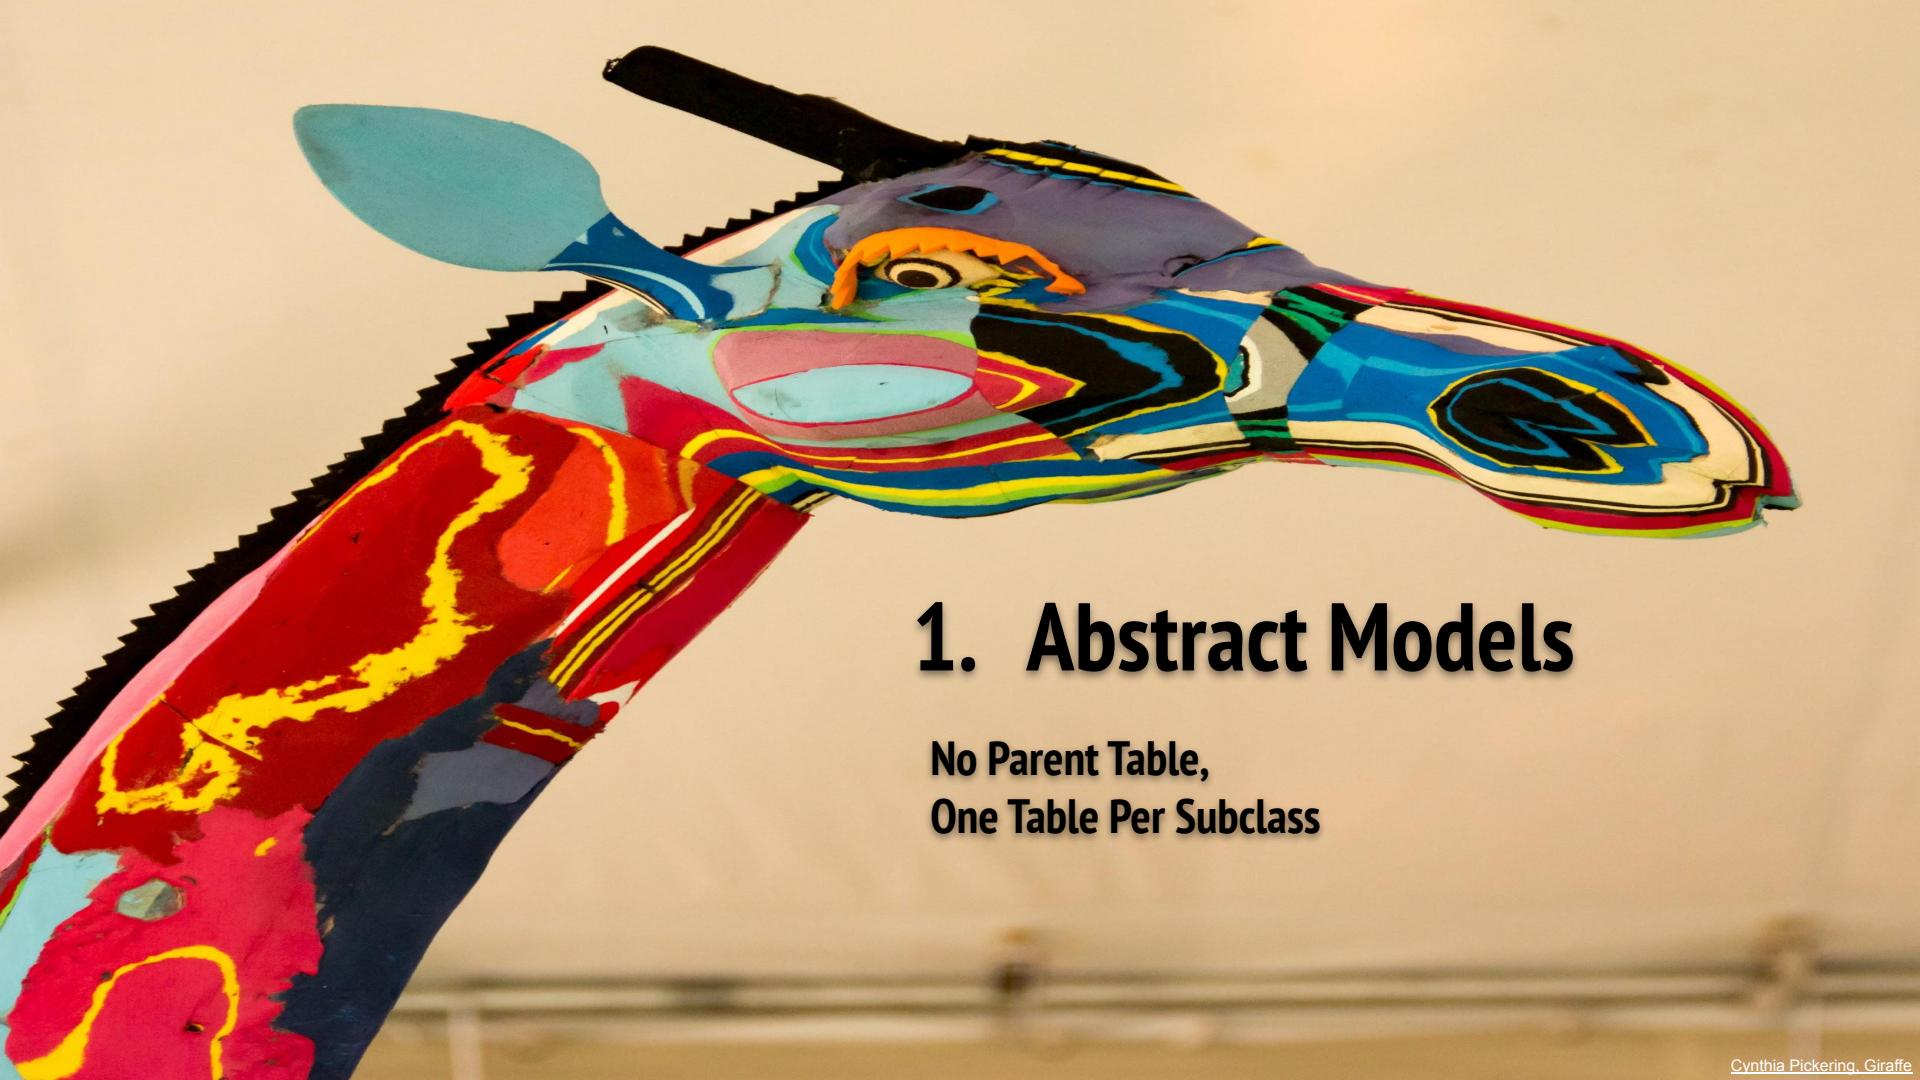
class Tongue(models.Model):
    color = models.CharField(max_length=255)
    giraffe = models.OneToOneField(
        Giraffe,
        on_delete=models.CASCADE,
    )
```

# Composition: OO UML vs DB Schema ERD



## 3 Types of Model Inheritance

Awkward!!!!

A large, colorful painting of a giraffe's head and neck, rendered in a style that obscures its form. The giraffe's body is primarily red with yellow highlights, while its head features a mix of blue, purple, orange, and black. It has a single large blue ear and a long, thin tongue. The background is a plain, light beige.

# 1. Abstract Models

**No Parent Table,  
One Table Per Subclass**

# Abstract Models: Base Definition

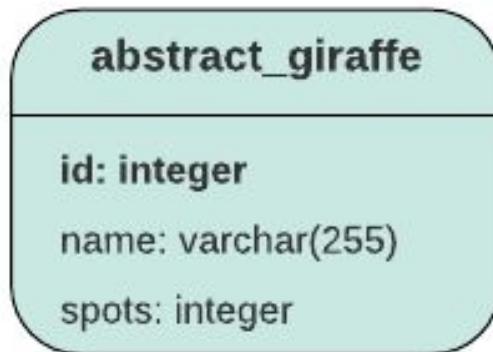
```
class Animal(models.Model):  
    class Meta:  
        abstract = True  
    name = models.CharField(max_length=255)  
  
    def speak(self):  
        return "##?!"
```

# Abstract Models: Derived Class Definition

```
class Giraffe(Animal):
    spots = models.IntegerField()

    def speak(self):
        return "Hum" # infrasonic
```

# Abstract Models: One Table Per Subclass



# Abstract Models: Query the SubClasses' table

```
In [1]: Giraffe.objects.first().speak()  
Out[1]: 'hum'
```

```
SELECT  
    "abstract_giraffe"."id",  
    "abstract_giraffe"."name",  
    "abstract_giraffe"."spots"  
FROM "abstract_giraffe"  
ORDER BY "abstract_giraffe"."id" ASC  
LIMIT 1
```

# Use Cases: Require Multiple Duplicated Fields

- BaseModel that is subclassed by many models
- TimeStampedModel overrides save to update the added and modified datetime fields



# Abstract Models: Ins and Outs

## Ins

- Reuses the parent class's fields and field related logic
- No performance overhead in querying one table

## Outs

- Parent class can not be used by itself
- Related models need multiple nullable foreign keys (Ex. zebra\_id and giraffe\_id are needed instead of a single animal\_id)

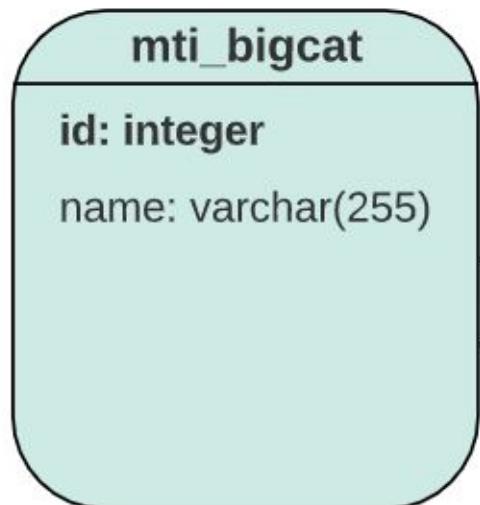
## 2. Multi-Table Inheritance (Concrete Inheritance)



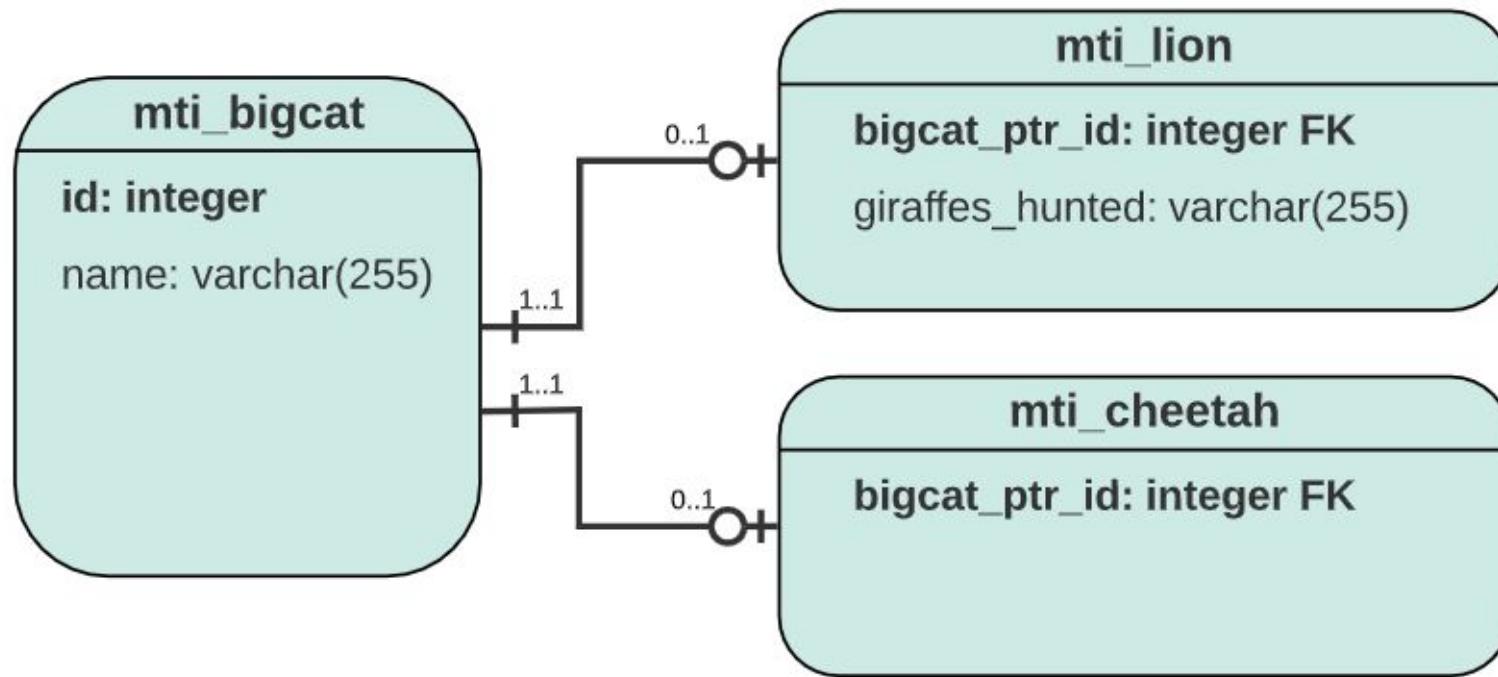
# MTI: No Meta Class Definition Needed

```
class BigCat(models.Model):  
    name = models.CharField(max_length=255)  
  
class Lion(BigCat):  
    giraffes_hunted = models.IntegerField()  
  
    def speak(self):  
        return "roar"
```

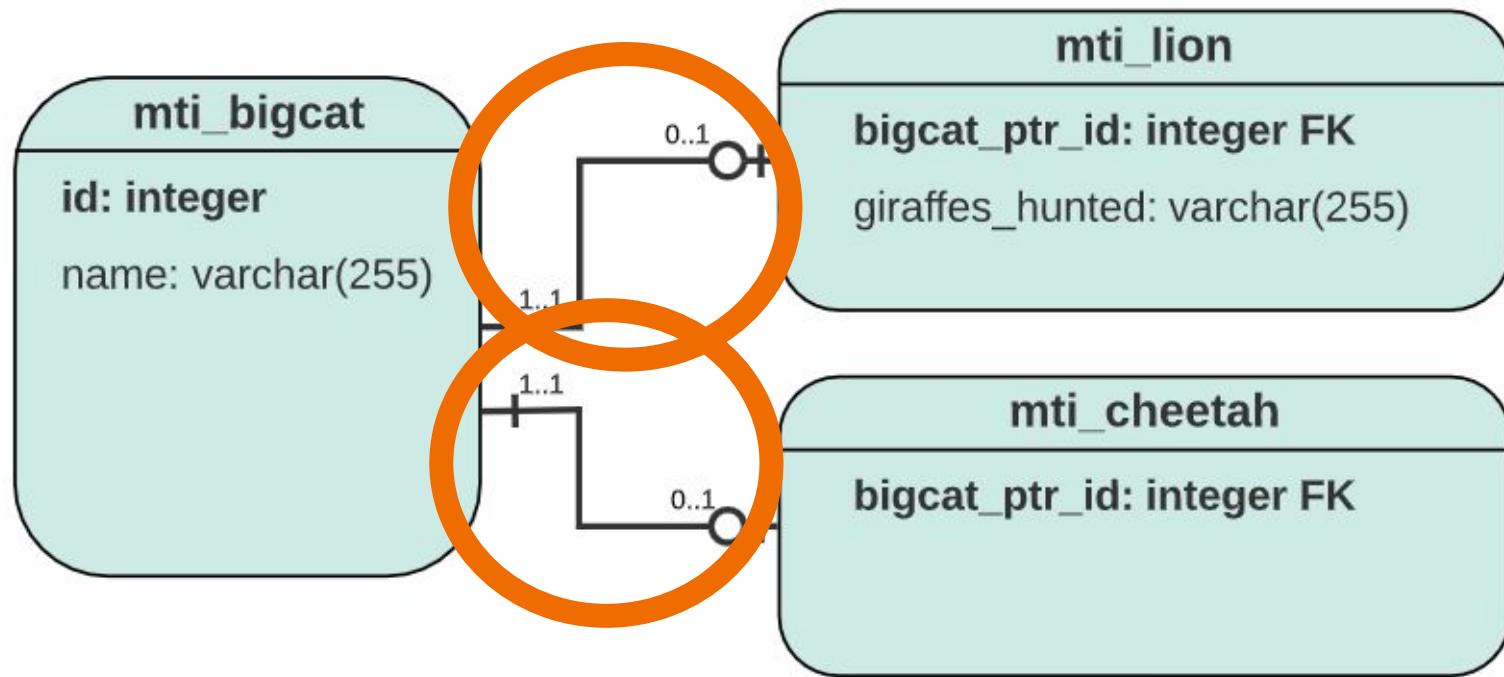
# MTI: Parent is concrete



# MTI: Parent is concrete; children use parent ptr as PK



# MTI: You could implement this explicitly with 1:1s



# MTI: Child is Queried with the Parent

```
In [1]: lions = list(Lion.objects.all())
SELECT
"mti_bigcat"."id",
"mti_bigcat"."name",
"mti_lion"."bigcat_ptr_id",
"mti_lion"."giraffes_hunted"
FROM "mti_lion"
INNER JOIN "mti_bigcat" ON (
"mti_lion"."bigcat_ptr_id" = "mti_bigcat"."id"
)
```

# MTI: Accessing the Parent Requires No Queries

```
In [2]: lions[0].bigcat_ptr # no Query, parent instance  
Out[2]: <BigCat: Simba (1)>
```

```
In [3]: lions[0].name          # no Query, parent fields aliased  
Out[3]: 'Simba'
```

```
In [4]: lions[0].speak()  
Out[4]: 'roar'
```

# MTI: Parent Query Starts Out Simple

```
In [5]: bigcats = list(BigCat.objects.all())
```

SELECT

```
"mti_bigcat"."id",  
"mti_bigcat"."name"  
FROM "mti_bigcat"
```

# MTI: Accessing each child executes a query

```
In [6]: bigcats[0].cheetah.speak()
```

```
SELECT
```

```
"mti_bigcat"."id",
"mti_bigcat"."name",
"mti_cheetah"."bigcat_ptr_id",
FROM "mti_cheetah"
INNER JOIN "mti_bigcat"
    ON ("mti_cheetah"."bigcat_ptr_id" = "mti_bigcat"."id")
WHERE "mti_cheetah"."bigcat_ptr_id" = 1
```

**RelatedObjectDoesNotExist: BigCat has no cheetah.**

# MTI: A query is needed for each subclass

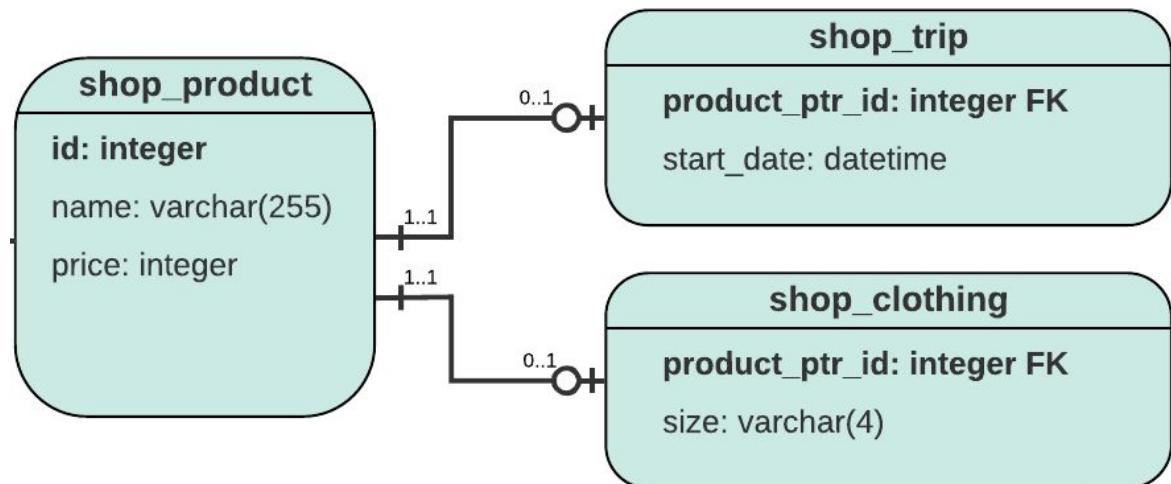
```
In [7]: bigcats[0].lion.speak()  
Out[7]: 'roar'
```

```
SELECT  
  "mti_bigcat"."id",  
  "mti_bigcat"."name",  
  "mti_lion"."bigcat_ptr_id",  
  "mti_lion"."giraffes_hunted"  
FROM "mti_lion"  
INNER JOIN "mti_bigcat" ON  
  ...
```

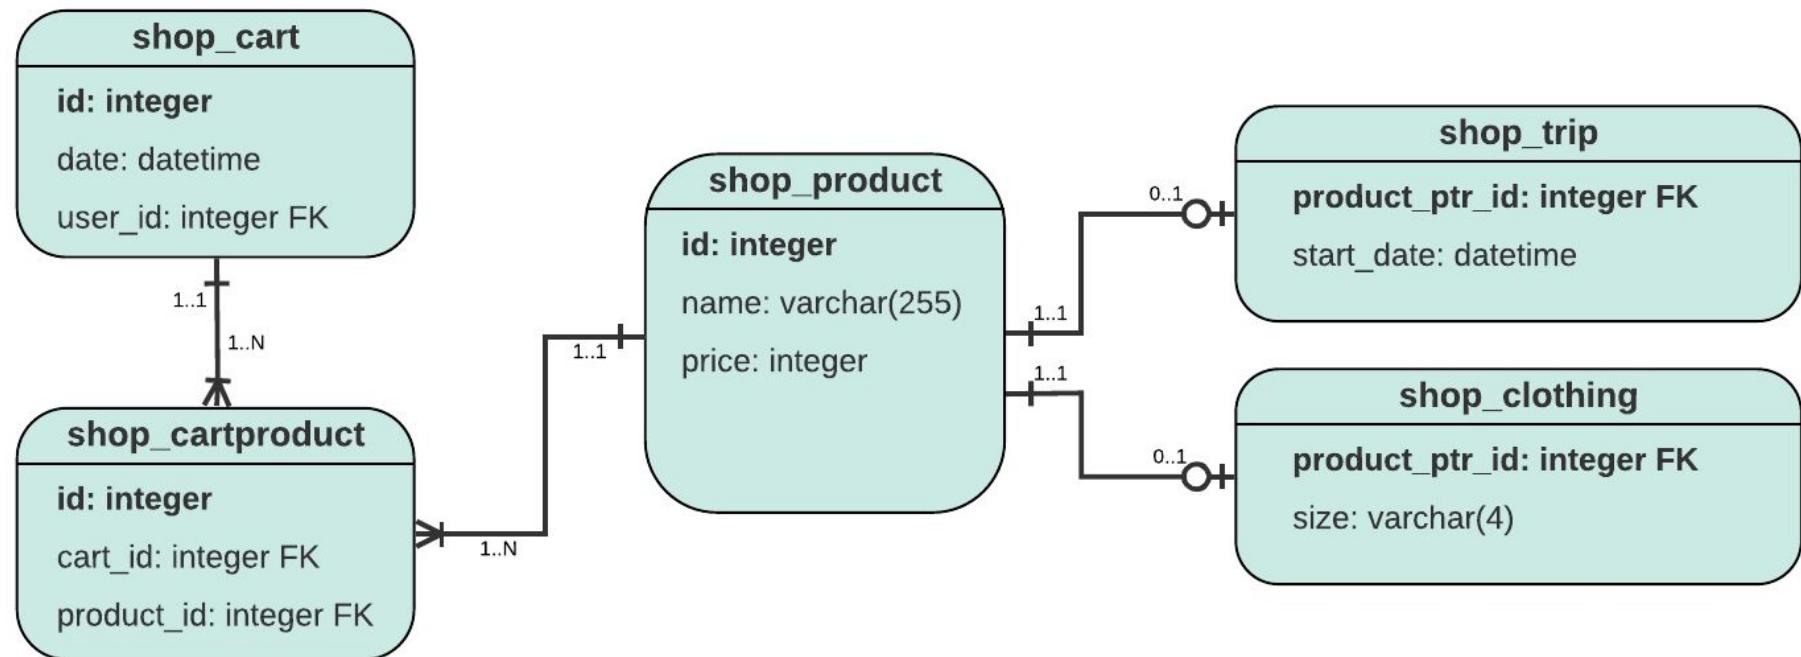
# Eager Load Optimization

1 prefetched query or LEFT JOIN per Subclass

# MTI Use Case: Shopping Cart



# MTI Use Case: Shopping Cart



## MTI Ins:

Common parent attributes can be queried together

## MTI: Outs

Inefficient queries hurt performance

Scaling is difficult

Lack of understanding



# 3. Proxy Models:

## One table, different behavior



# Proxy Models: Parent Class Definition

```
class Person(models.Model):  
    name = models.CharField(max_length=255)  
    person_type = models.CharField(max_length=3)  
  
    def characteristic(self):  
        return "Middle Earth Dweller"
```

# Proxy Models: Subclass Meta Class Sets Proxy

```
class Hobbit(Person):
    class Meta:
        proxy = True

    def characteristic(self):
        return "Hairy Feet"
```

# Proxy Models: 1 Table

proxy_person
<b>id: integer</b>
name: varchar(255)
person_type: varchar(3)



# Proxy Models: same data, different behavior

```
In [1]: Hobbit.objects.get(name='Frodo').characteristic()  
Out[1]: 'Hairy Feet'
```

```
In [2]: Person.objects.get(name='Frodo').characteristic()  
Out[2]: 'Middle Earth Dweller'
```

# Proxy Models: Subclass with Custom Manager

```
class Elf(Person):
    class Meta:
        proxy = True

    objects = ElfManager()

    def characteristic(self):
        return "Immortal"
```

# Proxy tables: Custom Manager Filters Elves

```
class ElfManager(models.Manager):
    def get_queryset(self):
        return super(ElfManager, self).get_queryset(
            ).filter(person_type='e')

    def create(self, **kwargs):
        kwargs.update({'person_type': 'e'})
        return super(ElfManager, self).create(**kwargs)
```

# Proxy Models with Custom Managers

```
In [1]: Person.objects.all()
```

```
Out[1]: <QuerySet [<Person: Frodo (1)>, <Person: Legolas (2)>]>
```

```
In [2]: Elf.objects.all()
```

```
Out[2]: <QuerySet [<Elf: Legolas (2)>]>
```

# Proxy: Same table, different filters and sorting

```
In: [1] Elf.objects.all()
```

```
SELECT
    "proxy_person"."id",
    "proxy_person"."name"
    "proxy_person"."person_type"
FROM    "proxy_person"
WHERE   "proxy_person"."person_type" = 'e'
```

# Proxy Models: Ins and Outs

**Ins:** Easy to modify the behavior of the subclasses

**Outs:** Fields used by any subclasses must be defined for all

## Use Cases:

- `OrderedModel` to sort on a field
- `ActiveModel` to filter `active=True`
- `CustomUserModel` for special `User` behavior (but it might be better to create a `UserProfile` 1:1 relationship instead)

# Downcasting and Single Table Inheritance

# Querying with Person Returns Person Instances

```
In [1]: people = Person.objects.all()
```

```
In [2]: people[0]
```

```
Out[2]: <Person: Frodo (1)>
```

```
In [3]: people[0].characteristic()
```

```
Out[3]: "Middle Earth Dweller"
```

# With “Downcasting” it returns the subclass instances

```
In [1]: Person.objects.all()
```

```
Out[1]: <QuerySet [<Hobbit: Frodo (1)>, <Elf: Legolas (2)>]>
```

```
In [2]: frodo = Person.objects.get(name="Frodo")
```

```
In [3]: frodo.characteristic()
```

```
Out[3]: "Hairy Feet"
```

# Downcasting Packages and Implementation

## Single-Table Inheritance

- [django-typed-models](#)
- [Django STI on the Cheap](#)

## Multi-Table Inheritance (CAREFUL!)

- [django-polymorphic](#) Prefetches 1 query per derived class
- [django-model-utils](#) Adds 1 join (select\_related) per derived class

# Single Table Inheritance Ins and Outs

## Ins

- **Performance:** One table means one query
- Related models need only one foreign key (`person_id`)

## Outs

- **Clutter and Bloat:** custom subclass fields are represented with nullable or defaulted columns

# STI and MTI have Similar Use Cases

## Single Table Inheritance

- Subclassed models **share most fields** and relationships
- Need to query across all subclasses

## Multi Table Inheritance

- Subclassed models **have many different fields**
- There is a ton of data (manually shard it by using multiple tables)

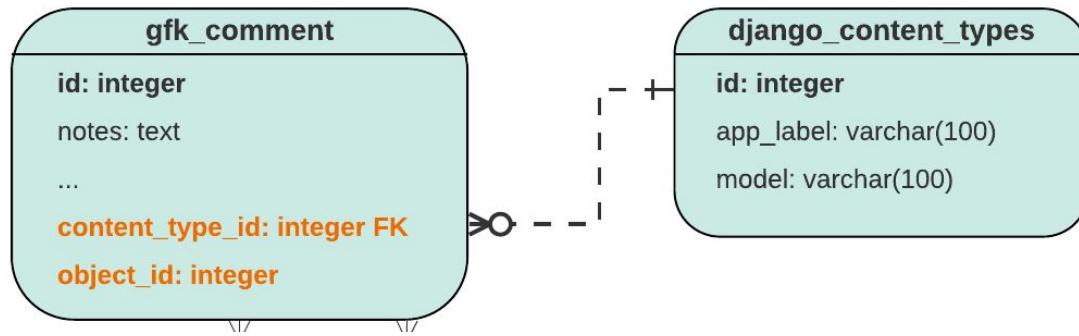
# The Best Type of Model Inheritance is No Inheritance

2 Alternative Approaches and  
Rethinking the Problem

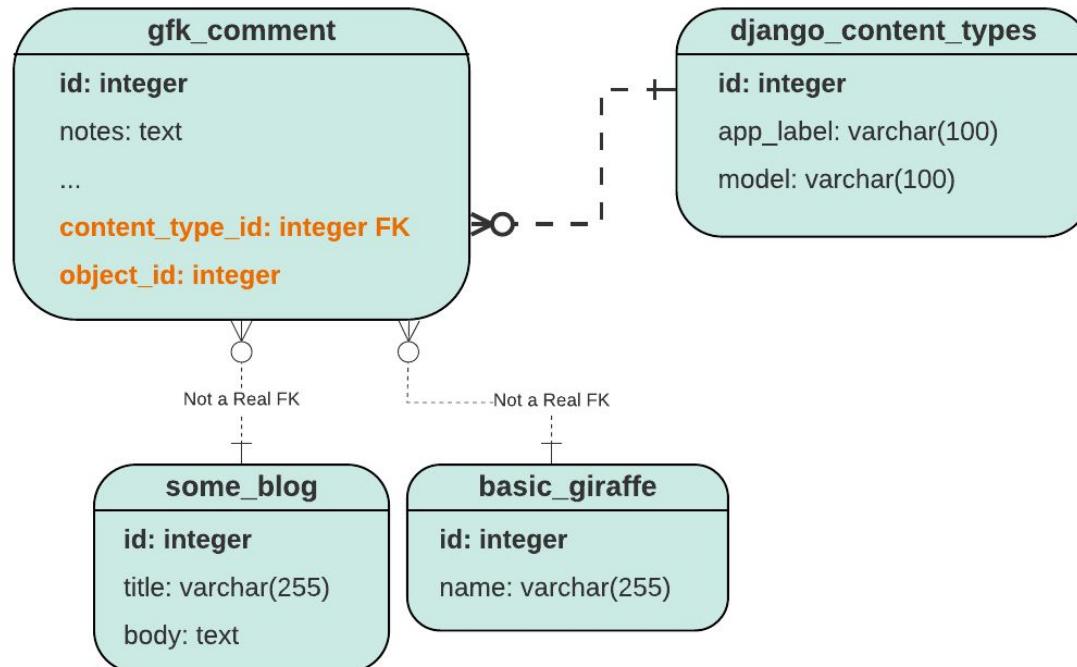
# 1. Polymorphism with Generic Foreign Keys



# Generic Foreign Keys: ERD Diagram



# Generic Foreign Keys: ERD Diagram



# Generic Foreign Keys: Ins and Outs

**Use Cases:** Tags, Comments, Likes, Favorites

**Ins:** Any model can be used and it does not need a migration

**Outs:**

- Code can become **hard to maintain**
  - dynamic type checking means the object could be anything  
select\_related is disallowed; custom sql needed for optimization
- **No DB referential integrity** could lead to dirty data
- **Performance hit** trying to query through to related objects

## 2. Unstructured Data ([JSONField](#) or [json-fields](#) Package)

**Use Cases:** Metadata, Events, Logs, Analytics

**Ins:**

- Avoids the clutter of nullable fields (vs STI)
- No need for related objects (vs MTI)

**Outs:**

- Tough to query against unstructured fields (except with Postgres)
- Validation and data integrity are not enforced by the database



Maybe we can  
rethink this

# Most is-a relationships can be expressed as has-a

Inheritance	Composition
A User is a Seller	A User has a SellerProfile
A Lion is a Carnivore	A Lion has a CarnivorousDiet
A Manager is an Employee	An Employee has a ManagerialJob

# Be Explicit: Your Future Self Thanks You

- Use Multiple foreign keys instead of inheritance
- Repeat yourself (a little bit) instead of **Abstract Models**
- **Implement Multi Table Inheritance explicitly** with foreign keys (ForeignKey and OneToOneField)

# Thank You!

<http://github.com/blythedunham/dmi>  
<http://linkedin/in/blythedunham>  
blythe@snowgiraffe.com



Blythe Dunham, Powpow for Breakfast